

# BVM 2019 / Tutorial: **Hands-On Deep Learning Using PyTorch (Part 1)**


**Christian Lucas**

Prof. Mattias P. Heinrich Group  
Institute of Medical Informatics  
University of Lübeck

# Convolutional Neural Networks (CNN)

- key idea: **learn convolutional kernels for features extraction**

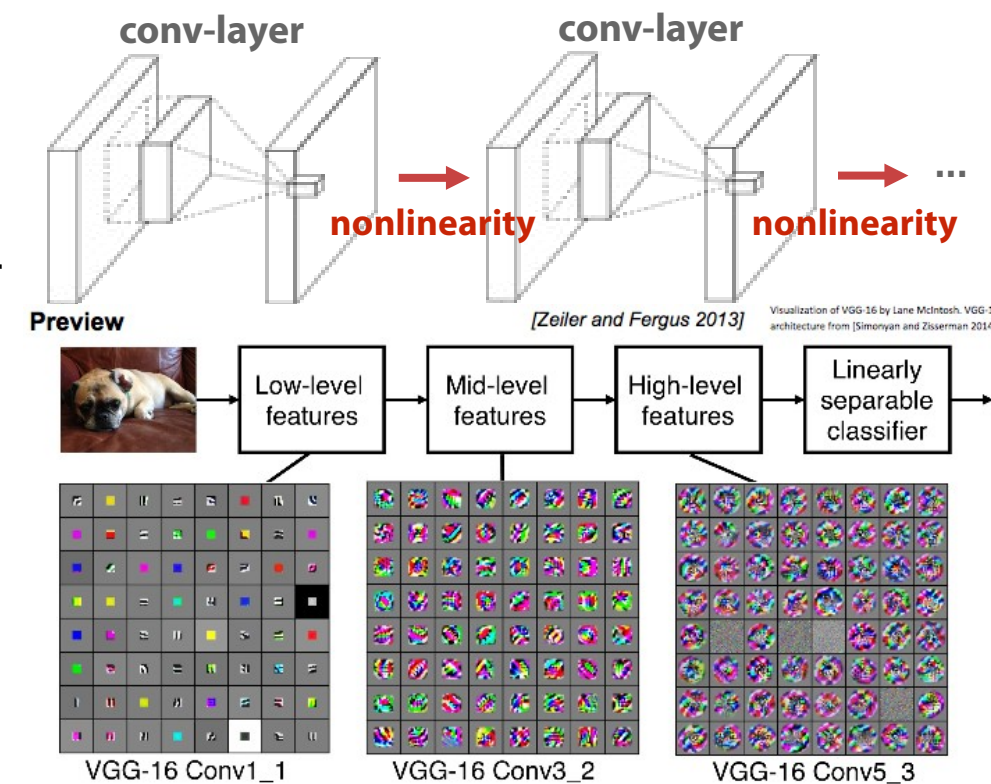

→ **use weight sharing**  
= **sparse connectivity**



# Convolutional Neural Networks (CNN)

- key idea: **learn convolutional kernels for features extraction**
- **Jointly learn everything within a deep architecture**  
→ Hierarchical feature transformations are jointly learned with classifier


→ **use weight sharing**  
**= sparse connectivity**



# Convolutional Neural Networks (CNN)

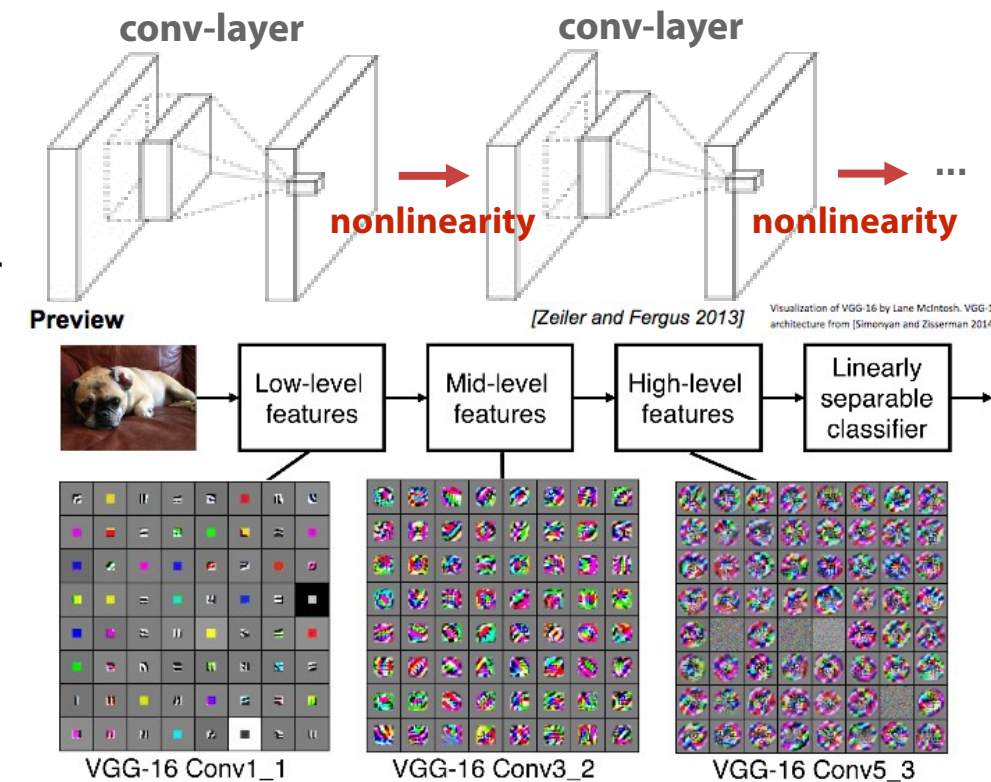

- key idea: **learn convolutional kernels for features extraction**
- **Jointly learn everything within a deep architecture**  
→ Hierarchical feature transformations are jointly learned with classifier

→ **use weight sharing**  
**= sparse connectivity**


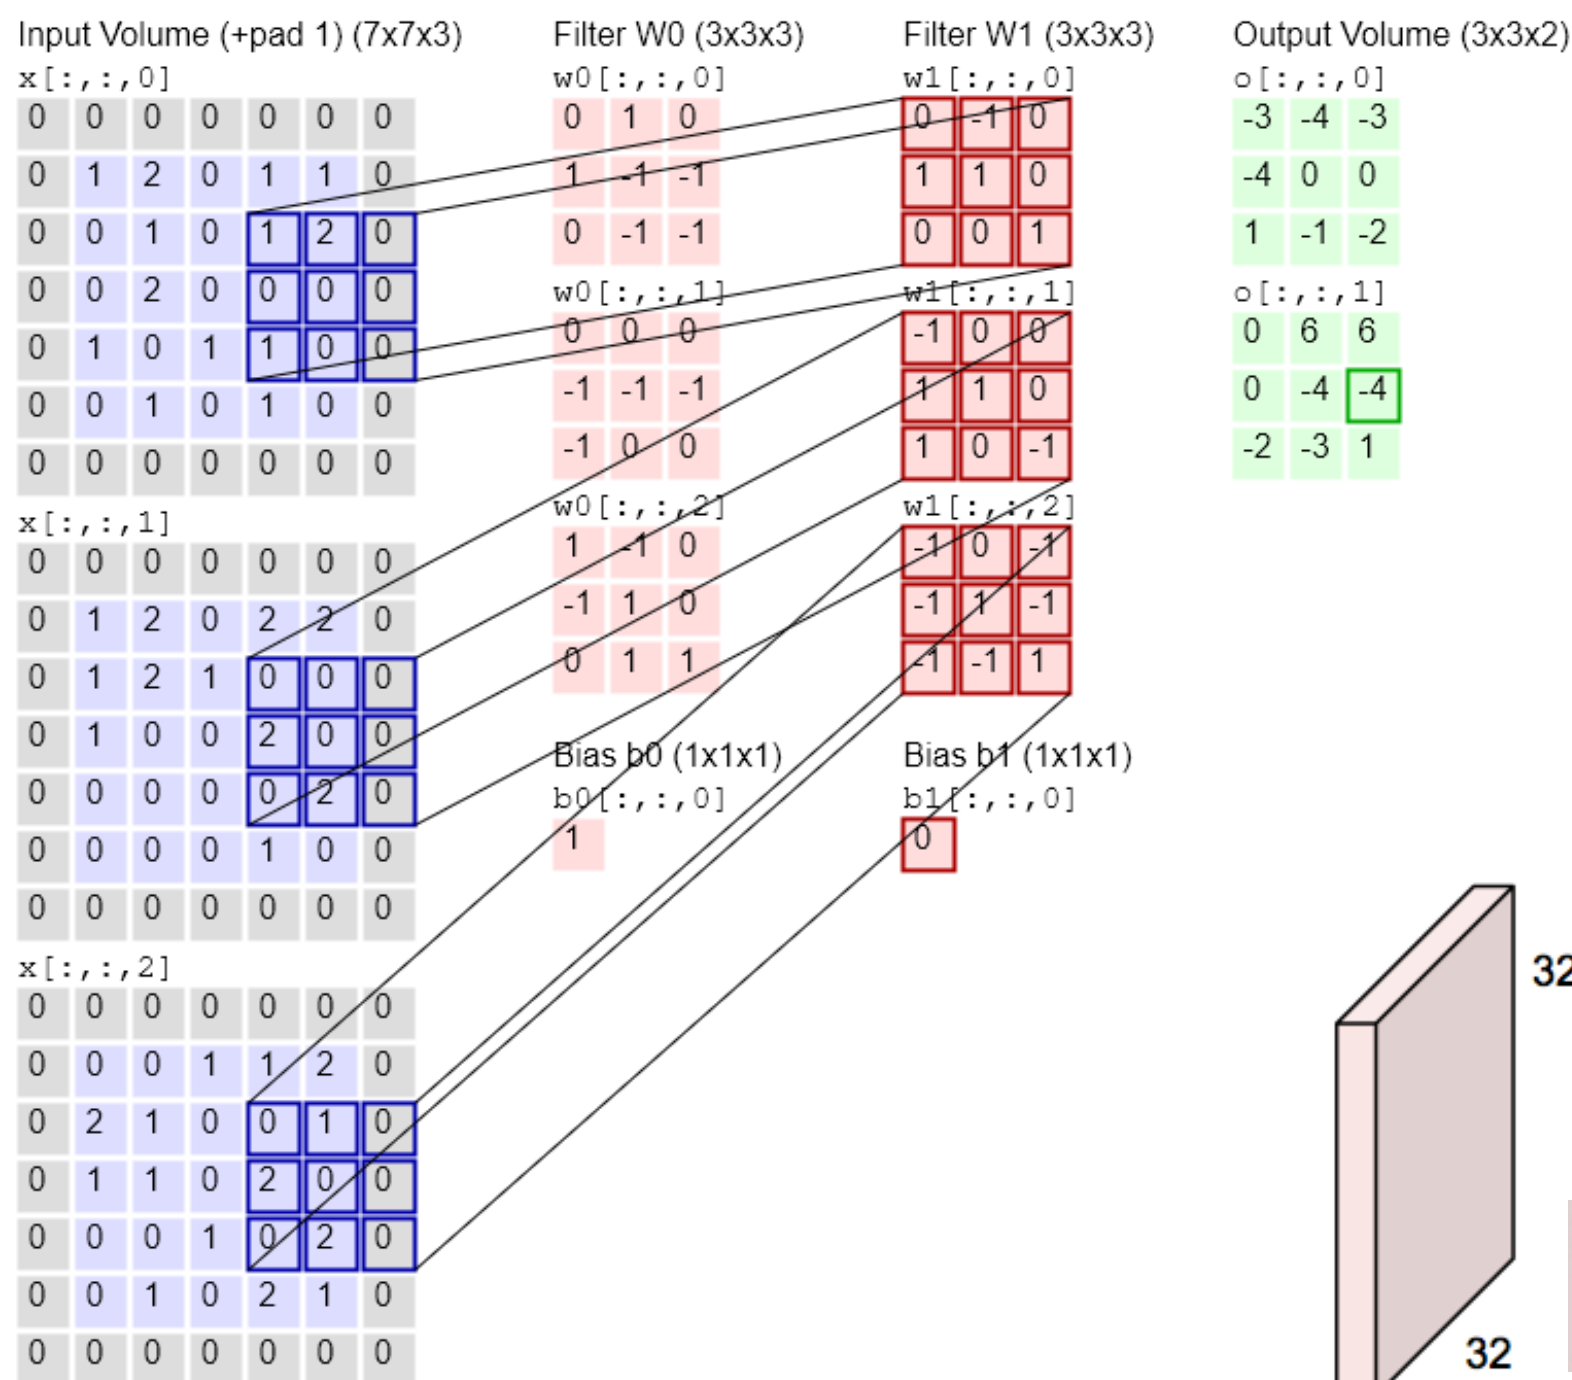


each filter is **applied to multi-channel patch** (e.g.  $8 \times 8 \times 3$ )


filter bank **generates multiple feature/activation maps**



# Convolutional Neural Networks (CNN)

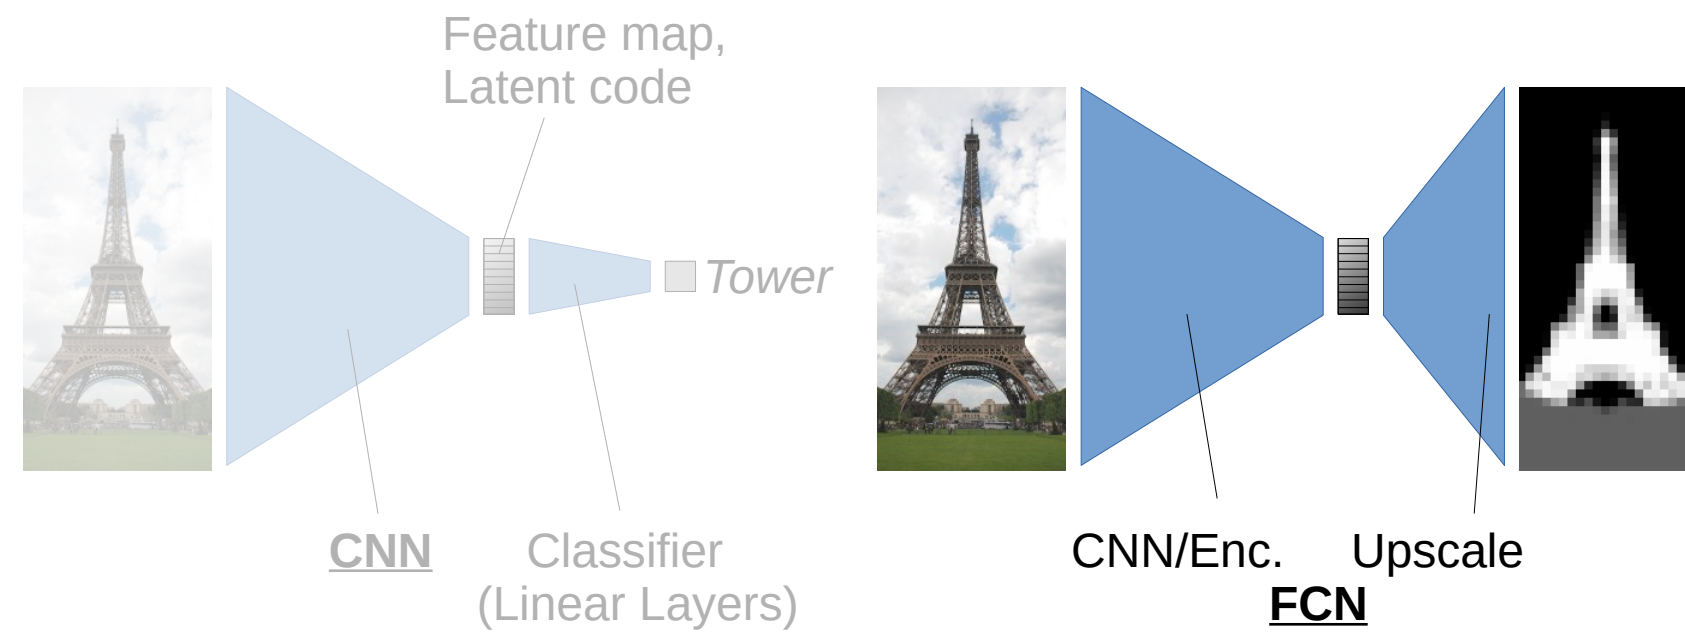


# Semantic Segmentation

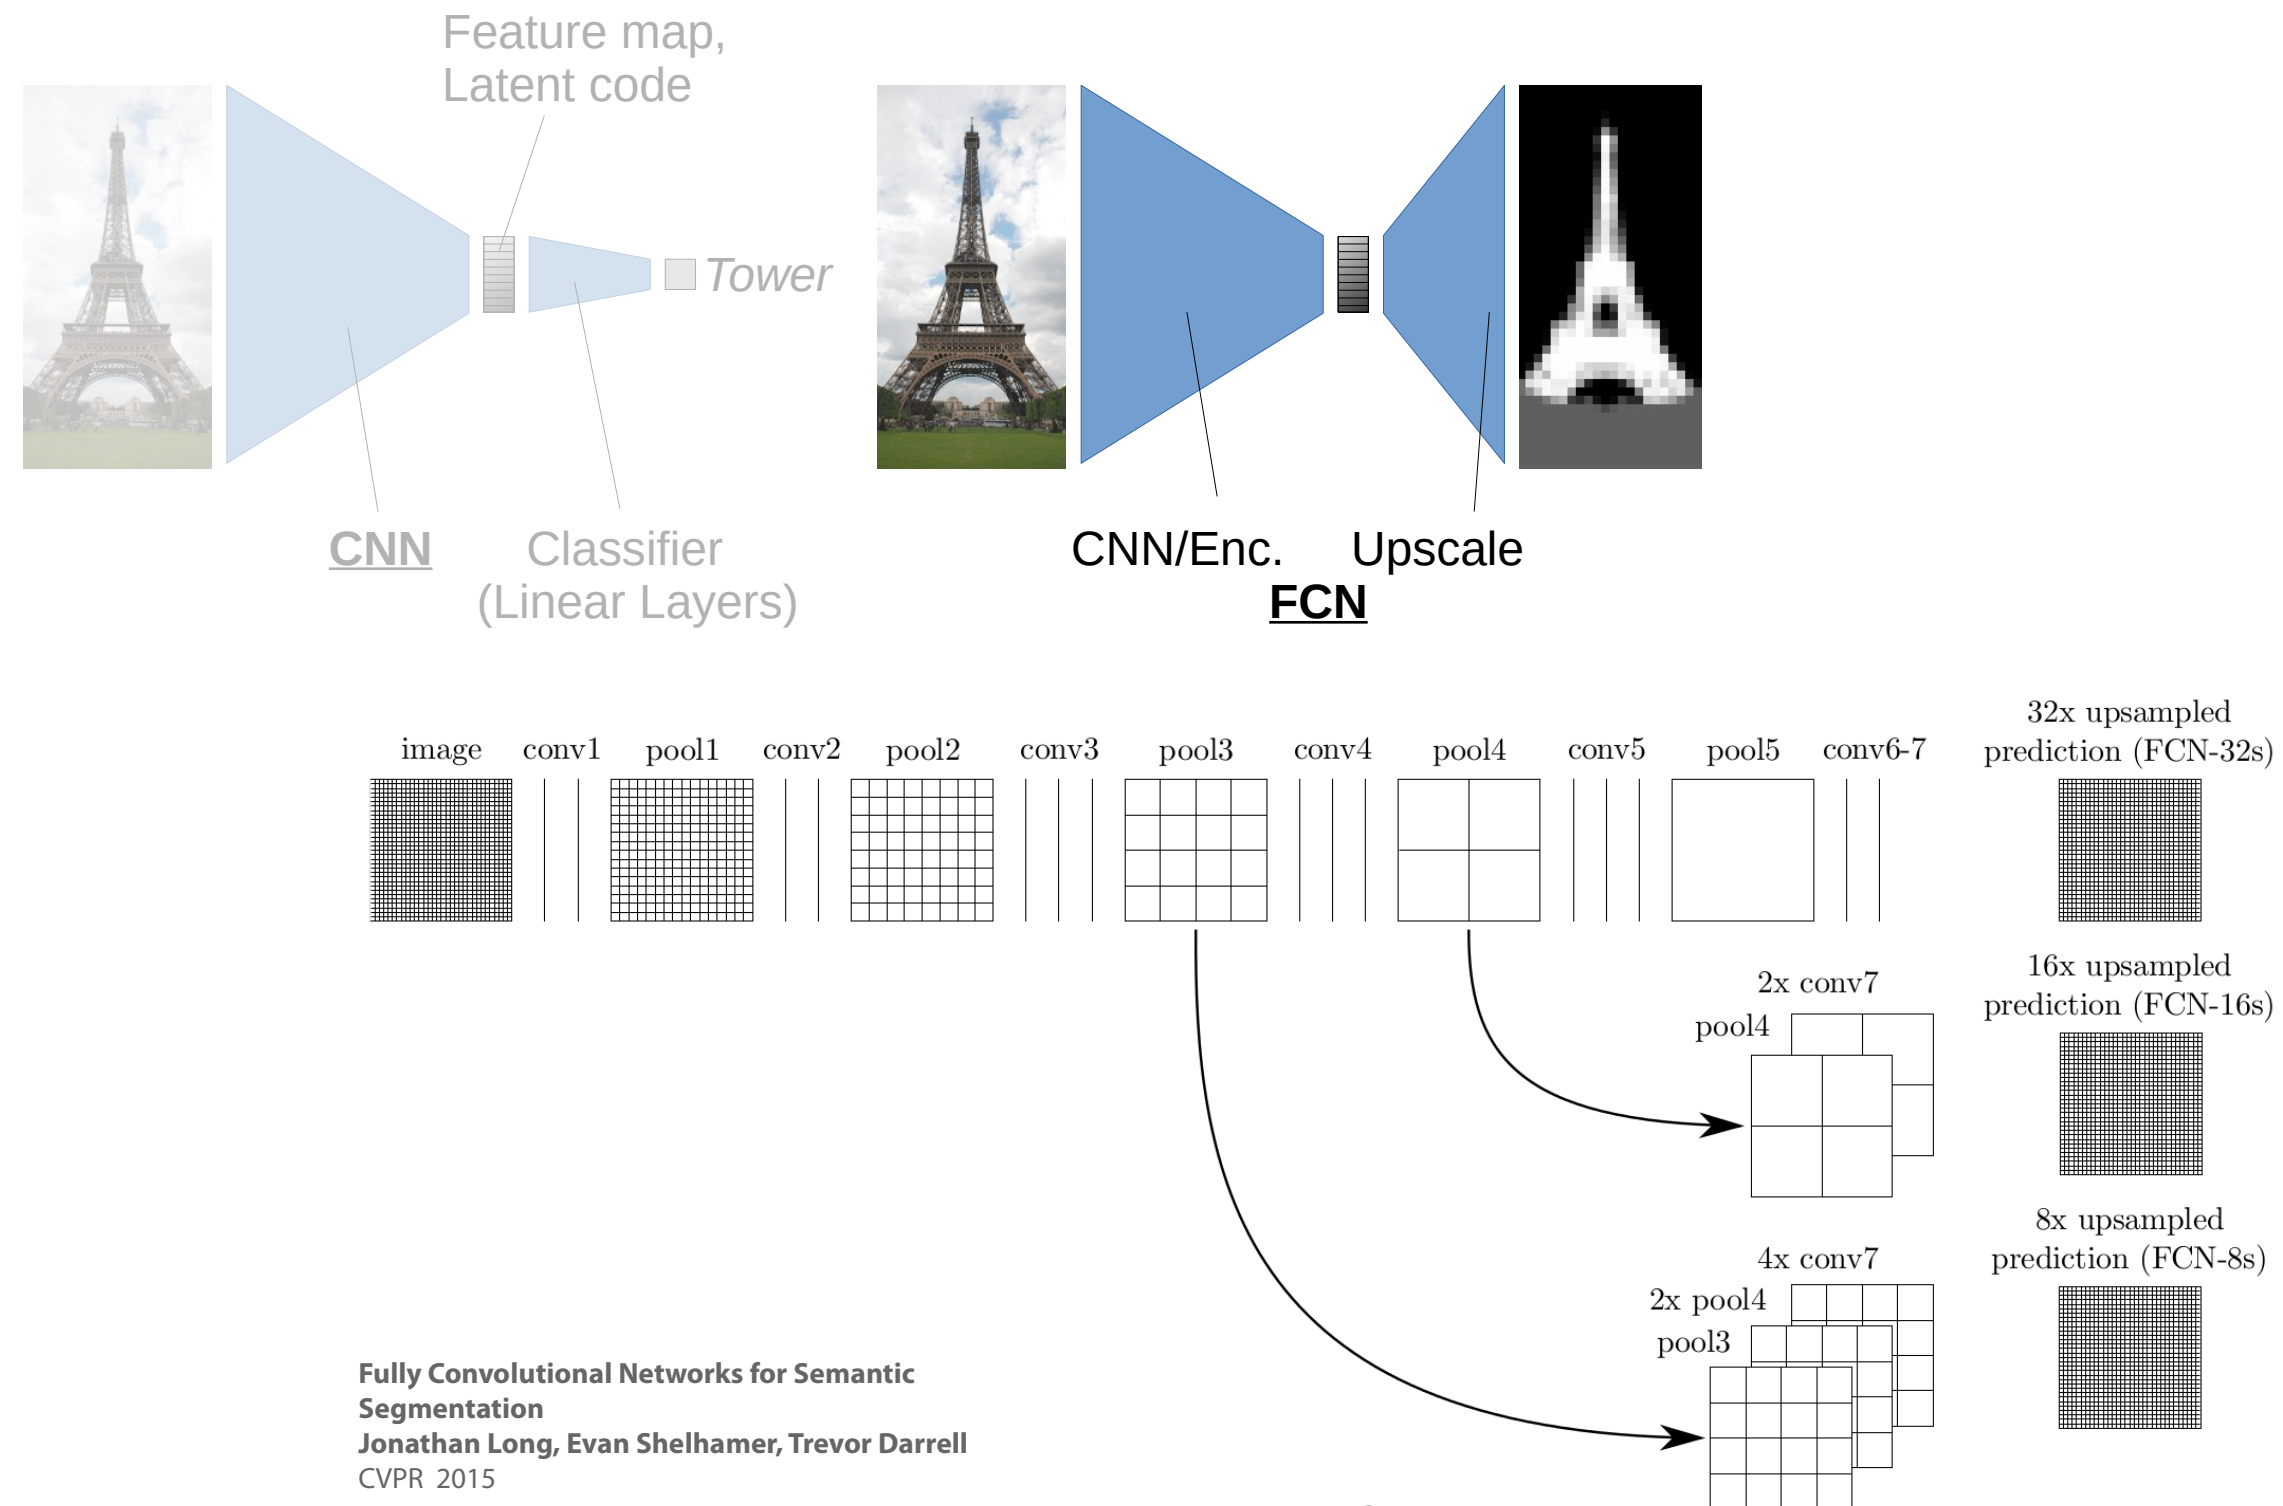




# Semantic Segmentation




# Semantic Segmentation







# Semantic Segmentation



# Semantic Segmentation

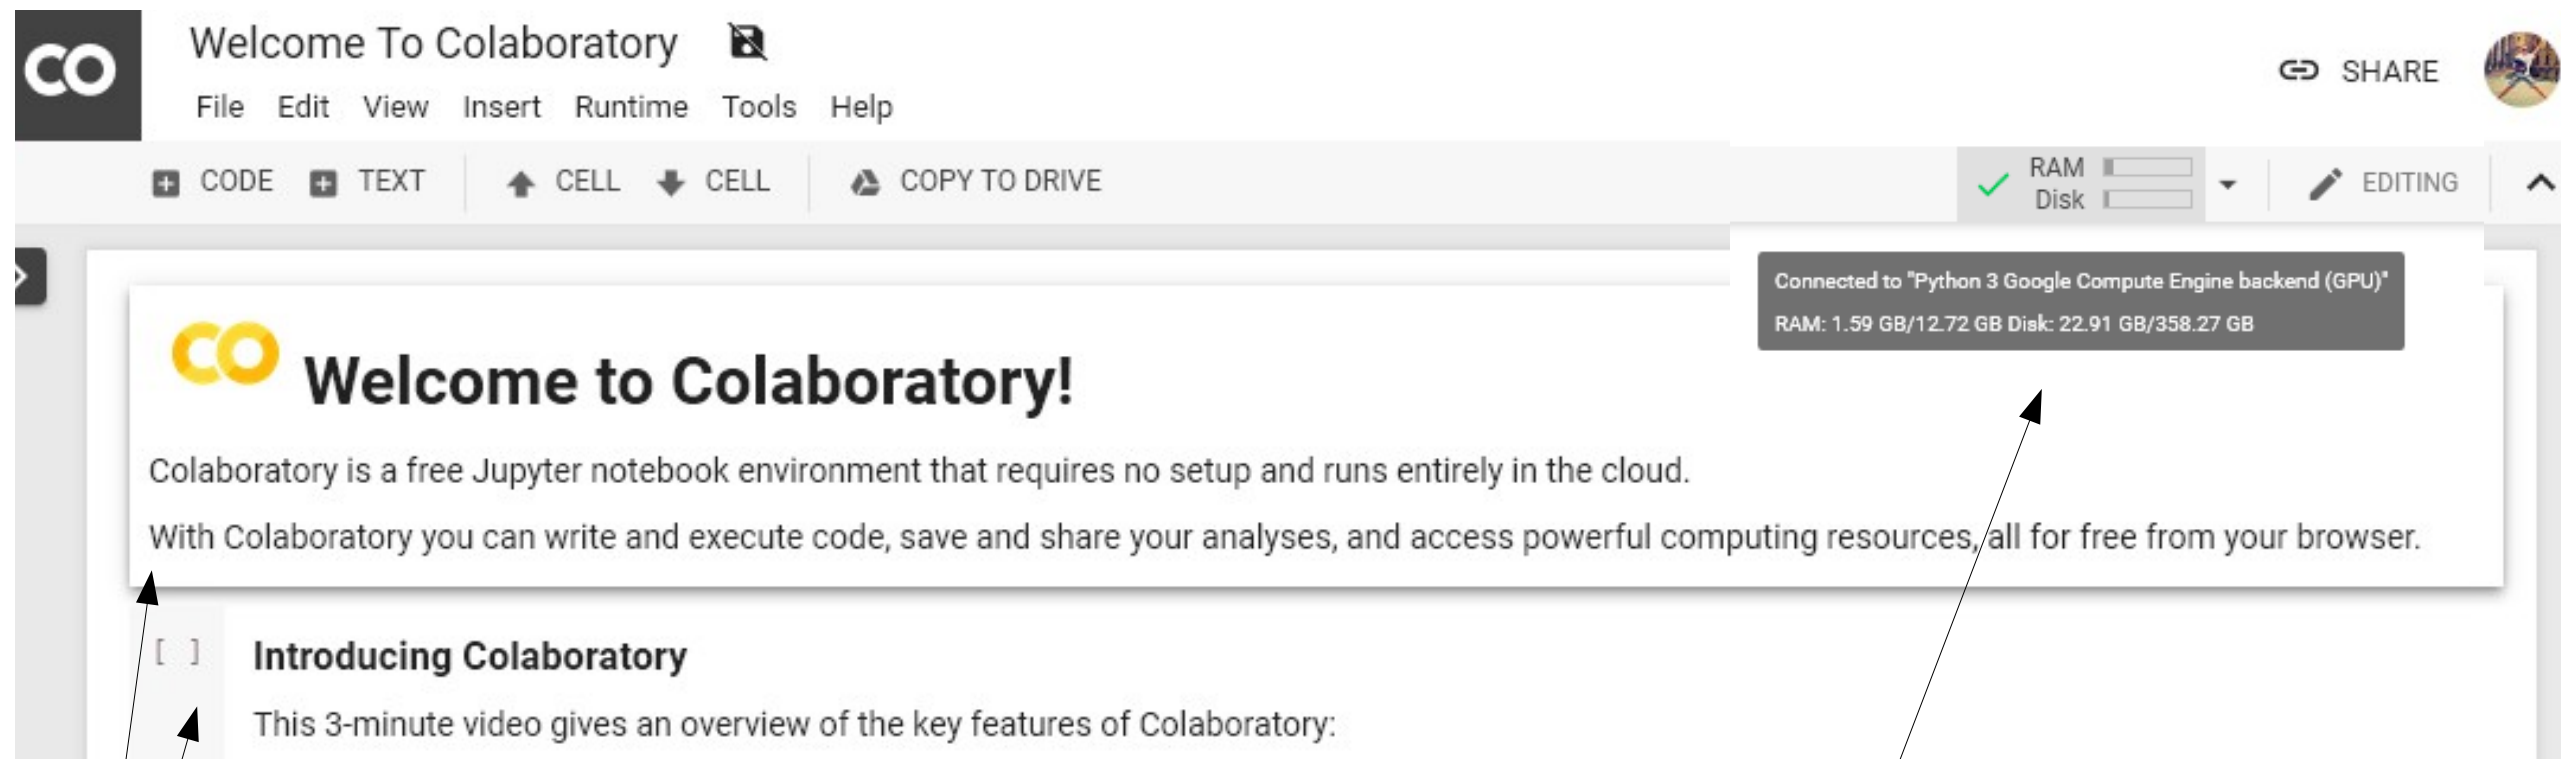


Swiss army knife for segmentation:  
Unified multi-resolution architecture



**U-Net: Convolutional Networks for Biomedical Image Segmentation**  
Olaf Ronneberger, Philipp Fischer, Thomas Brox  
MICCAI 2015, Lecture Notes in Computer Science  
(Vol. 9351), pp. 234-241, Springer

# Colab



Notebook cells


12GB GPU memory

`.cuda()`

- PyTorch 1.0 pre-installed
- Access to git: `!git clone https://github.com/name/repository.git`  
`import repository.pythonfile`
- Install other packages: `!pip install package` or `!setup.py install`


# Data

- Medical Segmentation Decathlon Task 02: **Heart data**  
([www.medicaldecathlon.com](http://www.medicaldecathlon.com))
- Segmentation of left atrium („linker Vorhof“)



Cardiac

**Target:** Left Atrium  
**Modality:** Mono-modal MRI  
**Size:** 30 3D volumes (20 Training + 10 Testing)  
**Source:** King's College London  
**Challenge:** Small training dataset with large variability



# Data Loading

```
class MyDataset(Dataset):
    def __init__(self, root_dir, transform=None):
        self.root_dir = root_dir
        self.transform = transform

    def __len__(self):
        return len(self.root_dir)

    def __getitem__(self, idx):
        img_name = os.path.join(self.root_dir, idx, '.nii.gz')
        image = io.imread(img_name)
        sample = {'image': image, 'index': idx}

        if self.transform:
            sample = self.transform(sample)

        return sample
```

`torch.utils.data.Dataset`: class representing a dataset  
    `__len__` returns the size of the dataset for `len(dataset)`  
    `__getitem__` for indexing (`dataset[i]` for  $i^{\text{th}}$  sample)

- Memory efficient implementation *on-demand* in `__getitem__`
- **Samples** as a dict `{'image': image, 'label': label}`
- Argument **transform** for any required pre-processing



# Data Loading

```
class MyDataset(Dataset):
    def __init__(self, root_dir, transform=None):
        self.root_dir = root_dir
        self.transform = transform

    def __len__(self):
        return len(self.root_dir)

    def __getitem__(self, idx):
        img_name = os.path.join(self.root_dir, idx, '.nii.gz')
        image = io.imread(img_name)
        sample = {'image': image, 'index': idx}

        if self.transform:
            sample = self.transform(sample)

        return sample
```

`torch.utils.data.Dataset`: class representing a dataset  
`__len__` returns the size of the dataset for `len(dataset)`  
`__getitem__` for indexing (`dataset[i]` for  $i^{\text{th}}$  sample)

- Memory efficient implementation *on-demand* in `__getitem__`
- **Samples** as a dict `{'image': image, 'label': label}`
- Argument `transform` for any required pre-processing

•  
•  
•

- Sampler for subsets (e.g., training set and validation set)

```
transformed_dataset = MyDataset(root_dir='/data/',
                                transform=transforms.Compose([
                                    Rescale(256),
                                    ToTensor()
                                ]))

dataloader = DataLoader(transformed_dataset,
                        batch_size=4,
                        shuffle=True,
                        num_workers=4)

for i_batch, sample_batched in enumerate(dataloader):
    # do whatever you'd like to do
```

```
class ToTensor(object):
    def __call__(self, sample):
        image, idx = sample['image'], sample['index']
        image = image.transpose((2, 0, 1))
        return {'image': torch.from_numpy(image), 'index': idx}
```

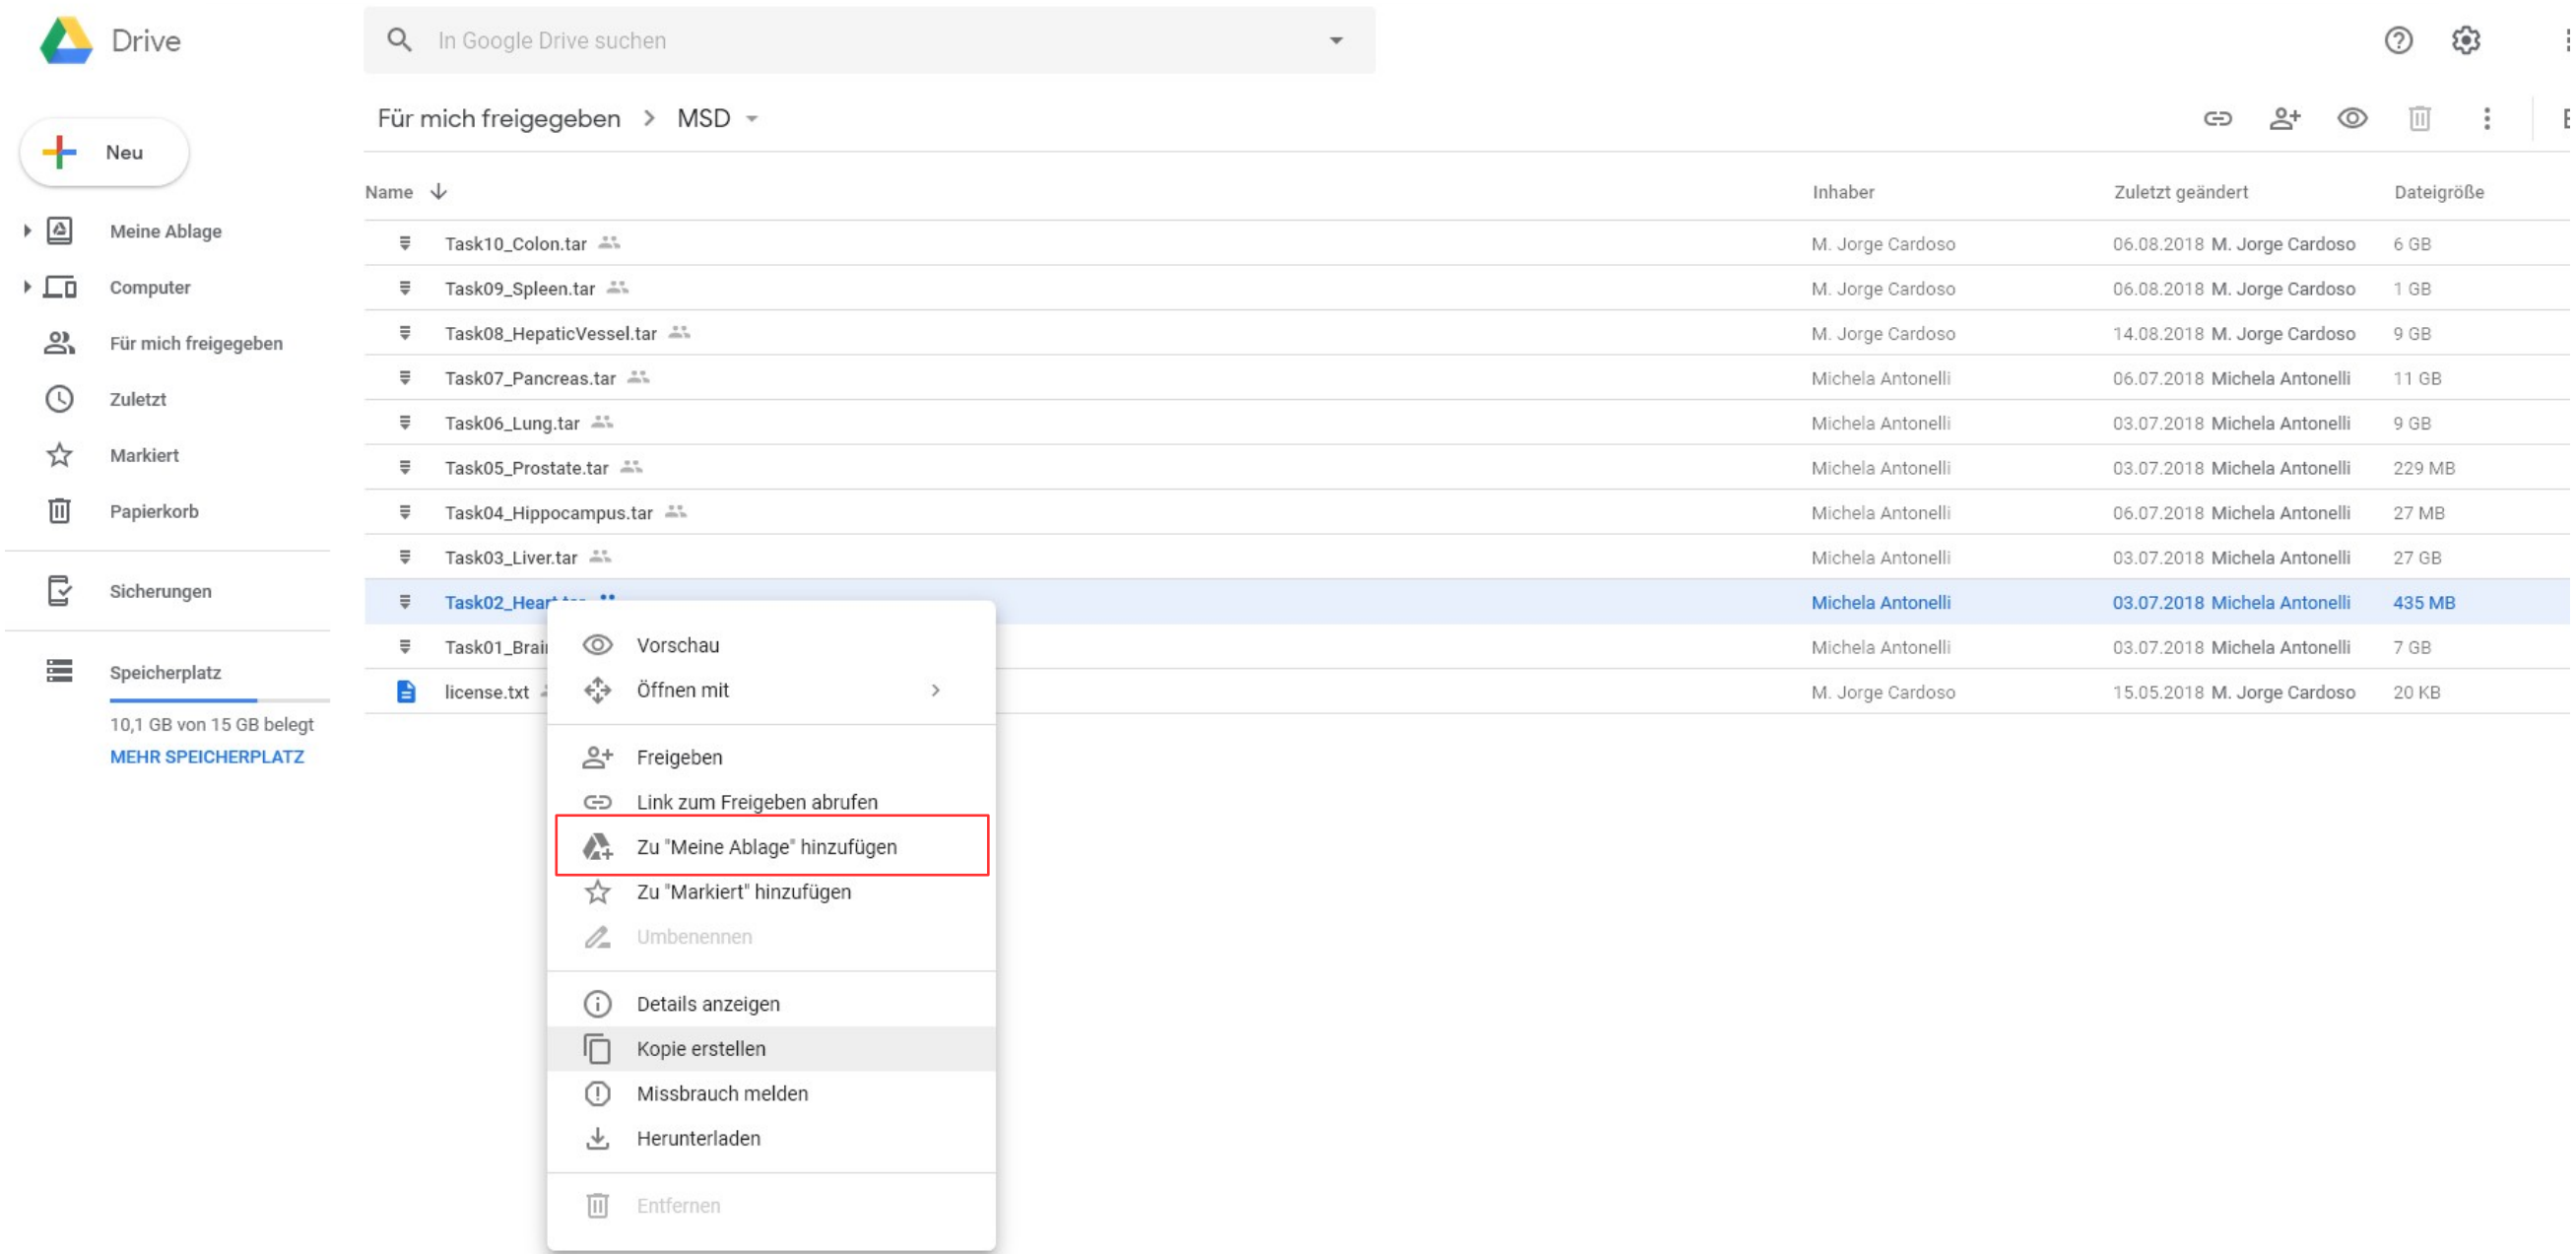
`torch.utils.data.DataLoader` is an iterator which provides:

- Batching the data
- Shuffling the data
- Load the data in parallel using multiprocessing workers.



# Hands on!

- Go to MSD drive and add **Task02\_Heart.tar** to your drive



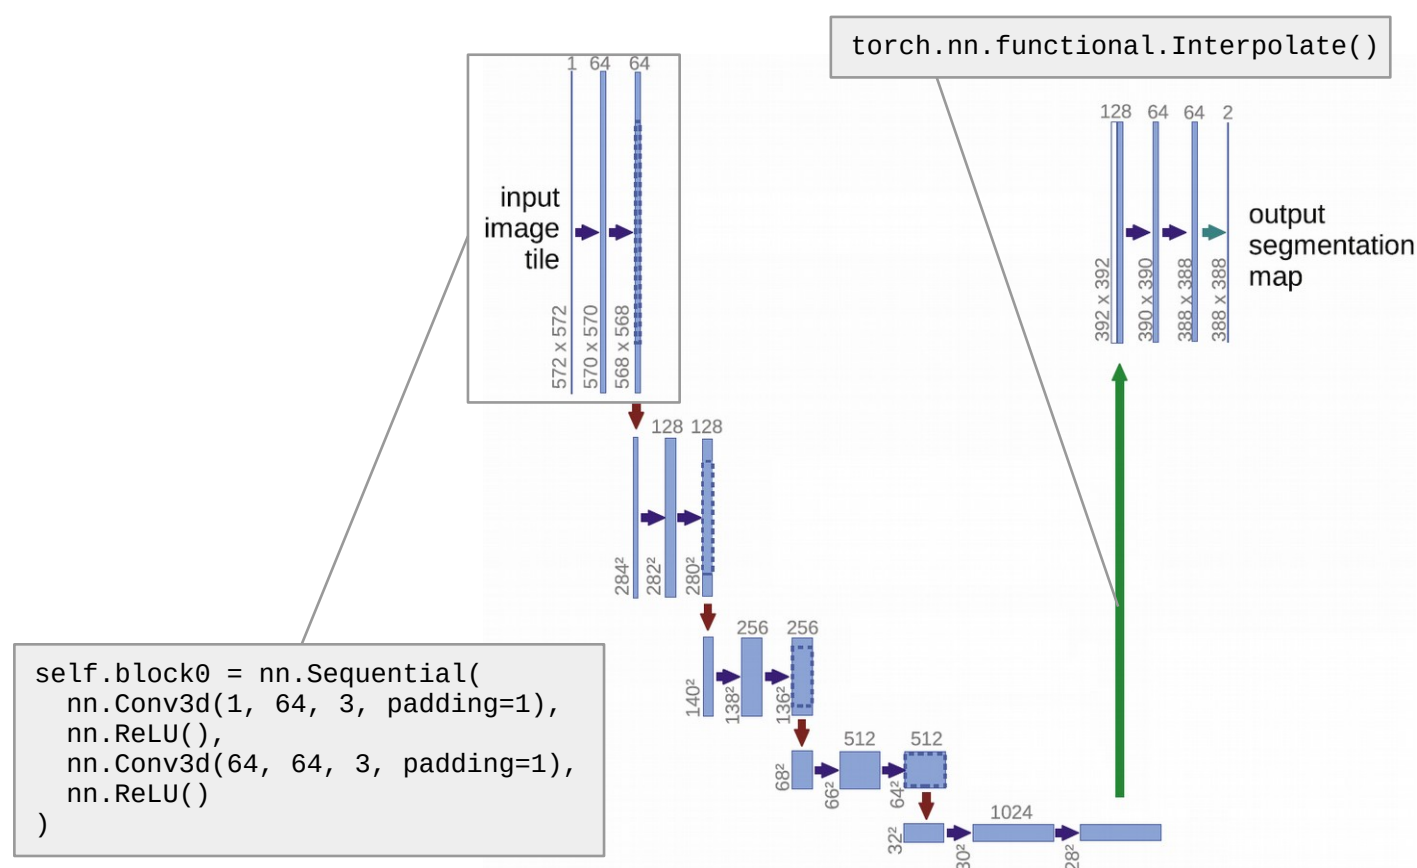
The screenshot shows the Google Drive interface. On the left, the sidebar displays navigation options: 'Meine Ablage', 'Computer', 'Für mich freigegeben', 'Zuletzt', 'Markiert', 'Papierkorb', 'Sicherungen', and 'Speicherplatz' (10,1 GB von 15 GB belegt). The main area shows the 'Für mich freigegeben' view of the 'MSD' folder. A table lists files with columns: Name, Inhaber, Zuletzt geändert, and Dateigröße. The file 'Task02\_Heart.tar' is highlighted. A context menu is open over this file, with the option 'Zu "Meine Ablage" hinzufügen' highlighted by a red rectangle.

Name	Inhaber	Zuletzt geändert	Dateigröße
Task10_Colon.tar	M. Jorge Cardoso	06.08.2018 M. Jorge Cardoso	6 GB
Task09_Spleen.tar	M. Jorge Cardoso	06.08.2018 M. Jorge Cardoso	1 GB
Task08_HepaticVessel.tar	M. Jorge Cardoso	14.08.2018 M. Jorge Cardoso	9 GB
Task07_Pancreas.tar	Michela Antonelli	06.07.2018 Michela Antonelli	11 GB
Task06_Lung.tar	Michela Antonelli	03.07.2018 Michela Antonelli	9 GB
Task05_Prostate.tar	Michela Antonelli	03.07.2018 Michela Antonelli	229 MB
Task04_Hippocampus.tar	Michela Antonelli	06.07.2018 Michela Antonelli	27 MB
Task03_Liver.tar	Michela Antonelli	03.07.2018 Michela Antonelli	27 GB
Task02_Heart.tar	Michela Antonelli	03.07.2018 Michela Antonelli	435 MB
Task01_Brain.tar	Michela Antonelli	03.07.2018 Michela Antonelli	7 GB
license.txt	M. Jorge Cardoso	15.05.2018 M. Jorge Cardoso	20 KB

- Open the first jupyter notebook in your `colab.research.google.com`


# Hands on!

- **Semantic segmentation** of the left atrium in the heart data via Deep Learning
- **Task 1:** Implement **Fully-Convolutional** Architecture
  - ✓ Double convolutional blocks
  - ✓ 2 poolings, and 1 final upsampling
  - ✓ Double the channel number, half the spatial resolution
  - ✓ To keep it easy: Use padding, so size of feature map after convolution is same as before
- **Task 2:** Extend FCN to **U-Net** (Symmetric decoder, more upsamplings, skip connections)



Task 1

16



Task 2

# Task 1: Solution

```
class FCN(nn.Module):
    def __init__(self):
        super().__init__()

        self.block0 = nn.Sequential(nn.BatchNorm3d(1),
                                     nn.Conv3d(1, 20, 3, padding=1),
                                     nn.ReLU(),
                                     nn.BatchNorm3d(20),
                                     nn.Conv3d(20, 20, 3, padding=1),
                                     nn.ReLU()
                                    )

        self.mp01 = nn.MaxPool3d(2, 2)

        self.block1 = nn.Sequential(nn.BatchNorm3d(20),
                                     nn.Conv3d(20, 40, 3, padding=1),
                                     nn.ReLU(),
                                     nn.BatchNorm3d(40),
                                     nn.Conv3d(40, 40, 3, padding=1),
                                     nn.ReLU()
                                    )

        self.mp12 = nn.MaxPool3d(2, 2)

        self.block2 = nn.Sequential(nn.BatchNorm3d(40),
                                     nn.Conv3d(40, 80, 3, padding=1),
                                     nn.ReLU(),
                                     nn.BatchNorm3d(80),
                                     nn.Conv3d(80, 80, 3, padding=1),
                                     nn.ReLU()
                                    )

        self.block3 = nn.Sequential(nn.Conv3d(80, 40, 1),
                                     nn.ReLU(),
                                     nn.Conv3d(40, 1, 1),
                                     )

        # 1 output channel segmentation

    def forward(self, inputs):

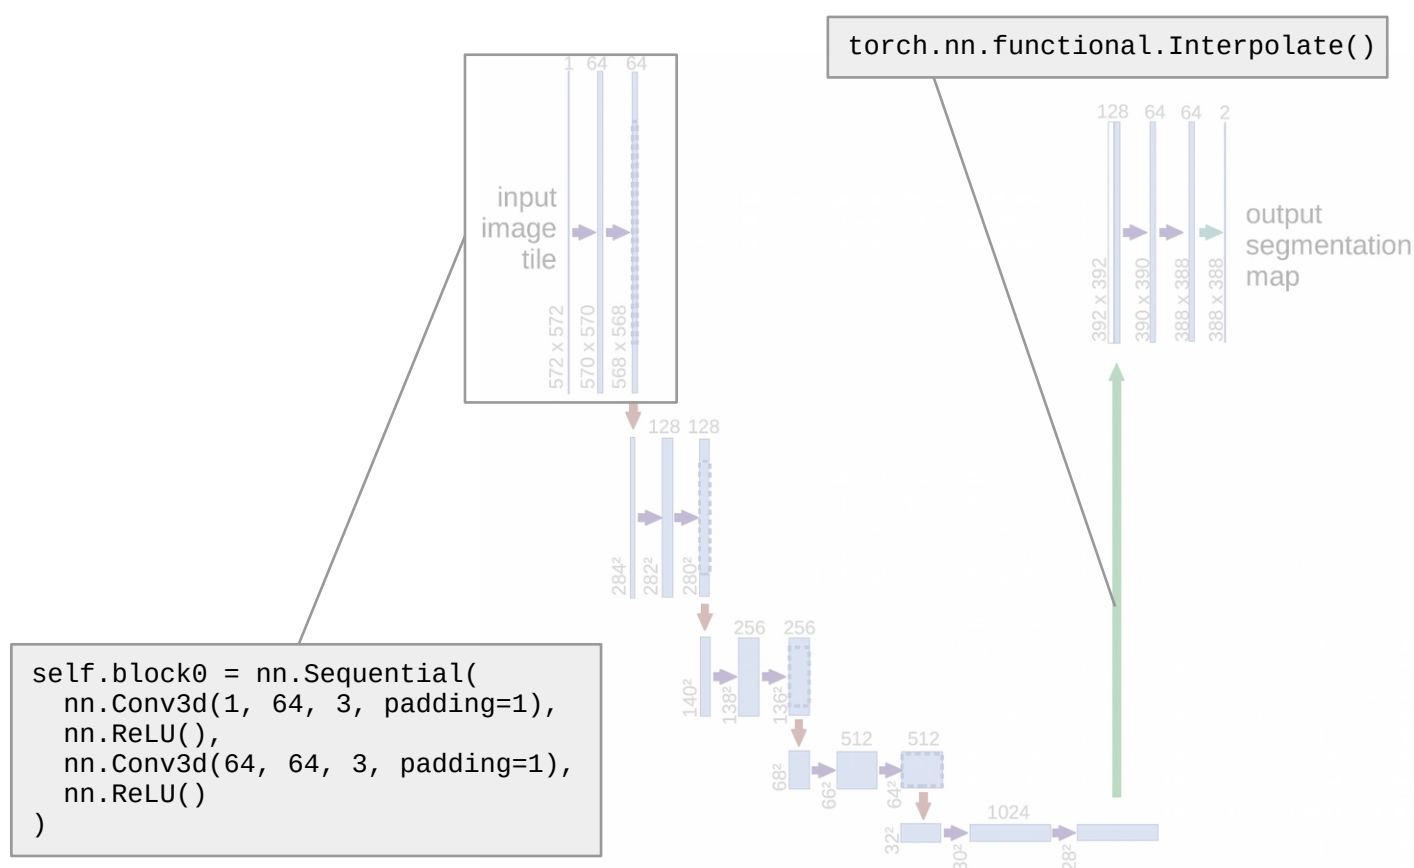
        output0 = self.block0(inputs)
        output1 = self.mp01(output0)
        output1 = self.block1(output1)
        output2 = self.mp12(output1)
        output2 = self.block2(output2)

        return self.block3(F.interpolate(output2, scale_factor=4))

        # 4x upsampling to recover from 2x 2-MaxPooling
```

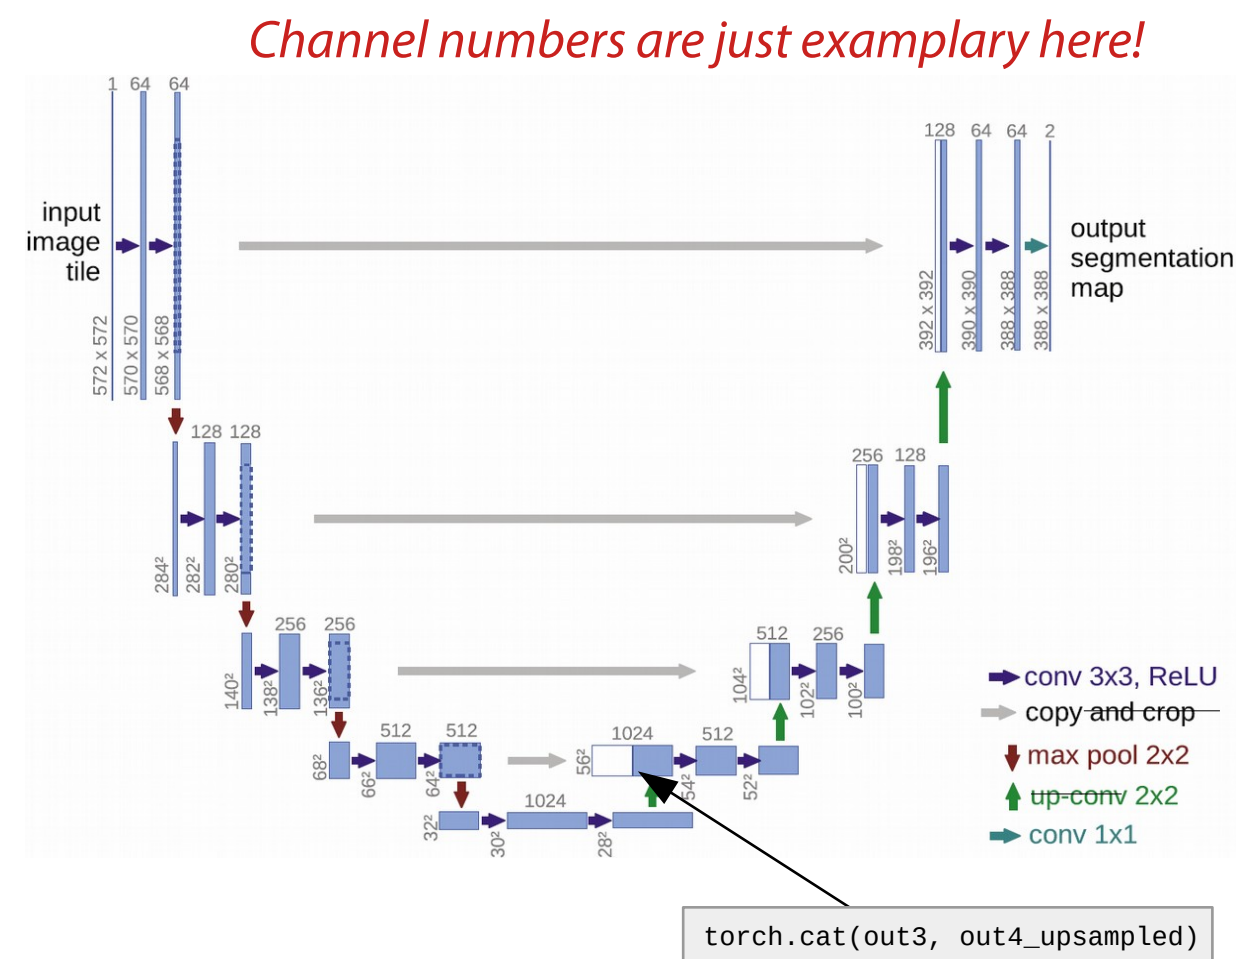
# Hands on!

- **Semantic segmentation** of the left atrium in the heart data via Deep Learning
- Task 1: Implement **Fully-Convolutional** Architecture
  - ✓ Double convolutional blocks
  - ✓ 2 poolings, and 1 final upsampling
  - ✓ Double the channel number, half the spatial resolution
  - ✓ To keep it easy: Use padding, so size of feature map after convolution is same as before
- Task 2: Extend FCN to **U-Net** (Symmetric decoder, more upsamplings, skip connections)



Task 1

18



Task 2

# Task 2: Solution

```
class Unet(nn.Module):
    def __init__(self):
        super().__init__()
        self.block0 = [...]
        self.mp01 = [...]
        self.block1 = [...]
        self.mp12 = [...]

        self.block2 = nn.Sequential(nn.BatchNorm3d(40),
                                    nn.Conv3d(40, 80, 3, padding=1),
                                    nn.ReLU(),
                                    nn.BatchNorm3d(80),
                                    nn.Conv3d(80, 80, 3, padding=1),
                                    nn.ReLU()
                                    )

        self.block3 = nn.Sequential(nn.BatchNorm3d(120),
                                    nn.Conv3d(120, 80, 3, padding=1),
                                    nn.ReLU(),
                                    nn.BatchNorm3d(80),
                                    nn.Conv3d(80, 40, 3, padding=1),
                                    nn.ReLU()
                                    )
        # 120 = upsample(block2) + skip(block1) = 80 + 40

        self.block4 = nn.Sequential(nn.BatchNorm3d(60),
                                    nn.Conv3d(60, 40, 3, padding=1),
                                    nn.ReLU(),
                                    nn.BatchNorm3d(40),
                                    nn.Conv3d(40, 20, 3, padding=1),
                                    nn.ReLU()
                                    )
        # 60 = upsample(block3) + skip(block0) = 40 + 20

        self.block5 = nn.Sequential(nn.Conv3d(20, 10, 1),
                                    nn.ReLU(),
                                    nn.Conv3d(10, 1, 1)
                                    )

    def forward(self, inputs):
        output0 = self.block0(inputs)
        output1 = self.mp01(output0)
        output1 = self.block1(output1)
        output2 = self.mp12(output1)
        output2 = self.block2(output2)

        output3 = F.interpolate(output2, scale_factor=2)
        output3 = self.block3(torch.cat([output3, output1], dim=1))
        output4 = F.interpolate(output3, scale_factor=2)
        output4 = self.block4(torch.cat([output4, output0], dim=1))

        return self.block5(output4)
```