

# Tutorial: Hands-On Deep Learning using pytorch (BVM 2019)

Mattias Heinrich, Christian Lucas, Max Blendowski

# Overview / learning outcomes

- 1) understand the principles of deep learning and its implementation with pytorch
- 2) build and train your own network and use autograd for iterative image registration (unsupervised)

Time	
14:00 - 14:40	Introduction to pytorch and image
14:40 - 15:45	<b>Practical Part I: build and train U-Net</b>
15:45 - 16:00	Coffee break
16:00 - 16:20	Intro to autograd and image registration
16:20 - 17:00	<b>Practical Part II: metrics + iterative alignment</b>
from 17:00	mingle with free BVM beer/wine

# Why pytorch is a good choice for deep learning:

"An open source deep learning platform that provides a seamless path from research prototyping to production deployment."



**originated from torch (lua-based) which offered:**

- a powerful N-dimensional array
- lots of routines for indexing, slicing, transposing, ...
- linear algebra routines
- neural network, and energy-based models
- numeric optimisation routines
- fast and efficient GPU support


Back-propagation  
uses the dynamically built graph

```
from torch.autograd import Variable

x = Variable(torch.randn(1, 10))
prev_h = Variable(torch.randn(1, 20))
W_h = Variable(torch.randn(20, 20))
W_x = Variable(torch.randn(20, 10))

i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
next_h = next_h.tanh()

next_h.backward(torch.ones(1, 20))
```



- great flexibility (dynamic graphs)
- very fast (C++/CUDA backends)
- immediate debugging python
- seamless installation (conda, ..)

Automatic differentiation in PyTorch

<https://openreview.net/pdf?id=BJJsrmfCZ>

# Medical Deep Learning using PyTorch

## Three main layers of the framework

**tensor**  
**data container / variable**  
 $B \times C \times H \times W \times D$   
*tracked by autograd*  
supports nearly all  
mathematical numpy  
functions (sum, exp, mm,  
etc.) and low-level  
operations: view, squeeze,  
broadcasting, slicing,  
typecasting

**nn**  
**neural network modules**  
with trainable parameters  
*interacts with optimiser*  
supports vast majority of  
(2D/3D) deep network  
operations: conv, batch-  
norm, relu, interpolate,  
pooling, etc.  
**nn.functional**  
similar functionality but  
without optimising  
parameters, + *grid\_sample*

**optim**  
**stochastic gradient  
descent**  
optimisers with momentum,  
e.g. *Adam*. Uses autograd to  
calculate gradient update  
steps w.r.t. to loss:  
`net.train()`  
`optimizer.zero_grad()`  
`output = net(input)`  
`loss = criterion(output,  
label)`  
`loss.backward()`  
`optimizer.step()`

**torchvision.datasets & .models**  
pre-processed and annotated computer vision datasets and pre-trained 2D models  
**torch.utils.data.Dataset & .checkpoint**  
abstract class for defining a dataset with batch-loader, checkpointing for saving memory



# A simplistic introduction

```
A = torch.Tensor([2,4])  
A.requires_grad = True  
B = 1/3*A**3  
B.sum().backward()  
print(A.grad)  
  
tensor([ 4., 16.] )
```

pytorch's most important functionality is **automatic differentiation**

# Some helpful pytorch commands

Create new **Tensors** with: `torch.zeros(2,3)` `randn()` `ones()` `linspace` `arange` etc  
use **broadcasting** and **views**:

```
import torch
```

```
x_grid = torch.linspace(-1,1,5)  
y_grid = torch.arange(0,7)  
xy_grid = x_grid.view(-1,1) + y_grid.float().view(1,-1)  
print(xy_grid)
```

```
tensor([[ -1.0000,  0.0000,  1.0000,  2.0000,  3.0000,  4.0000,  5.0000],  
        [ -0.5000,  0.5000,  1.5000,  2.5000,  3.5000,  4.5000,  5.5000],  
        [  0.0000,  1.0000,  2.0000,  3.0000,  4.0000,  5.0000,  6.0000],  
        [  0.5000,  1.5000,  2.5000,  3.5000,  4.5000,  5.5000,  6.5000],  
        [  1.0000,  2.0000,  3.0000,  4.0000,  5.0000,  6.0000,  7.0000]])
```

~~Docstring~~ can be displayed when cursor is on function (1. in line) with Shift+Tab (not in colab)

```
x_grid = torch.linspace(-1,1,5)  
y_grid = torch.arange(0,7)
```

Docstring:

```
linspace(start, end, steps=100, out=None, dtype=None, layout=torch.strided, device  
=None, requires_grad=False) -> Tensor
```

```
torch.nn.functional.gr  
torch.nn.functional.grad  
torch.nn.functional.grid_sample  
torch.nn.functional.group_norm
```

**Tab auto-extends function  
names**



# helpful python / pytorch commands

**Indentation** is always necessary in python (but no semi colons needed),

→ Tab or Shift-Tab with cursor at beginning of line,  
methods can be called from Tensor object with .operator()

```
with torch.no_grad():  
    print(xy_grid.sum())
```

```
tensor(105.)
```

Jupyter notebooks are based on cell execution (similar to %% in Matlab)

**Ctrl+Enter** or **Alt+Enter** runs code of current cell

(Alt+Enter in addition jumps into a new line)

**Funktionen** werden in python mit `def name(arguments, ..):` deklariert

Vorsicht keine klare Unterscheidung zwischen globalen und lokalen Variablen (Lesezugriff auf beides möglich)

```
import torch  
import torch.nn.functional as F  
import matplotlib.pyplot as plt  
data = torch.load('uebung6_data.pth')  
img_train = data['img_train']
```

```
def resample(img):  
    x_grid = torch.linspace(-1,1,img.size(2)//3).view(1,img.size(2)//3,1,1)\  
        .repeat(1,1,img.size(3)//3,1)  
    y_grid = torch.linspace(-1,1,img.size(3)//3).view(1,1,img.size(3)//3,1)\  
        .repeat(1,img.size(2)//3,1,1)  
    xy_grid = torch.cat((y_grid,x_grid),3)  
    img_resampled = F.grid_sample(img,xy_grid)  
    return img_resampled
```

```
img_resampled = resample(img_train[0:1,:,:,:])  
plt.imshow(img_resampled[0,0,:,:].detach(),'gray')  
plt.show()
```

tensor  
data container / variable

# list of widely used functions

see more details in documentation <https://pytorch.org/docs/stable/index.html>

**matrix multiply:** `torch.matmul(tensor1, tensor2)`

**clip** the values of a tensor: `torch.clamp(input, min, max)`

**swap** two specified dimensions: `torch.transpose(input, dim0, dim1)`

short hand for swapping Dim=0 and Dim=1 is `.t()`

for >3 dimensions: `.permute(*dims) : *dims (int...)` – the desired ordering of dimensions

**reshape** the size of dimensions, same memory (total number of elements must be equal): `.view(*shape)`

**select** with tensor of integer indices: `torch.index_select(input, dim, index)`

**slicing:** similar to MatLab, but declared as (start:end:stride) e.g. select every second row of 2D Tensors starting from 3:

`x_ = x[3::2,:]`

**repeat** a tensor along an axis: `expand(*sizes) : *sizes (torch.Size or int...)` – the desired expanded size

remove a (any) dimension of size 1: `.squeeze()`

add a dimension of size 1: `.unsqueeze(dim)` (useful for broadcasting)

**concatenate** list of tensors: `torch.cat(tensors, dim) : y = torch.cat((x,x),0)`

return number of dimensions `.dim()`; and **return their sizes** `.size(dim)`;

change datatype of tensor: `.type(dtype=None) : string` e.g. 'float' or simply `.float().cuda()`



nn  
neural network

# common deep learning functions

see more details in documentation <https://pytorch.org/docs/stable/nn.html>

nonlinearity ReLU (with extra argument for slope default=0):

```
torch.nn.LeakyReLU()
```

convert predictions into pseudo probabilities (for classification)

```
torch.nn.LogSoftmax(dim=None) including logarithm
```

```
torch.nn.Softmax(dim=None)
```

**convolutional layers** for learning features and pooling layers for smoothing / downsampling:

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)
```

```
torch.nn.AvgPool2d(kernel_size, stride=None, padding=0, ceil_mode=False, count_include_pad=True)
```

transposed convolution (used in backward path of Conv2d and classical U-Nets):

```
torch.nn.ConvTranspose2d(in_channels, out_channels, kernel_size, stride=1, padding=0, output_padding=0, groups=1, bias=True, dilation=1)
```

loss function for categorical classification / segmentation (use together with log\_softmax)

```
torch.nn.NLLLoss(weight=None, size_average=None, ignore_index=-100, reduce=None, reduction='mean')
```

and loss function for continuous regression (L1 norm, absolute differences):

```
torch.nn.L1Loss(size_average=None, reduce=None, reduction='mean')
```

extract all overlapping patches (n-D) e.g. to create sparse networks or implement own filters:


```
torch.nn.Unfold(kernel_size, dilation=1, padding=0, stride=1)
```

some words on the difference to nn.functiona

# Stochastic Gradient Descent (SGD)

conventional optimisation would sum gradient over all data points

**inefficient** for larger datasets, **cancelling out** of opposing gradients → **mini-batches**



find **solution** starting with random init of **w**  
**Gradient Descent:**  $\vec{w}^{(k)} = \vec{w}^{(k-1)} - \eta \frac{dL}{d\vec{w}}$

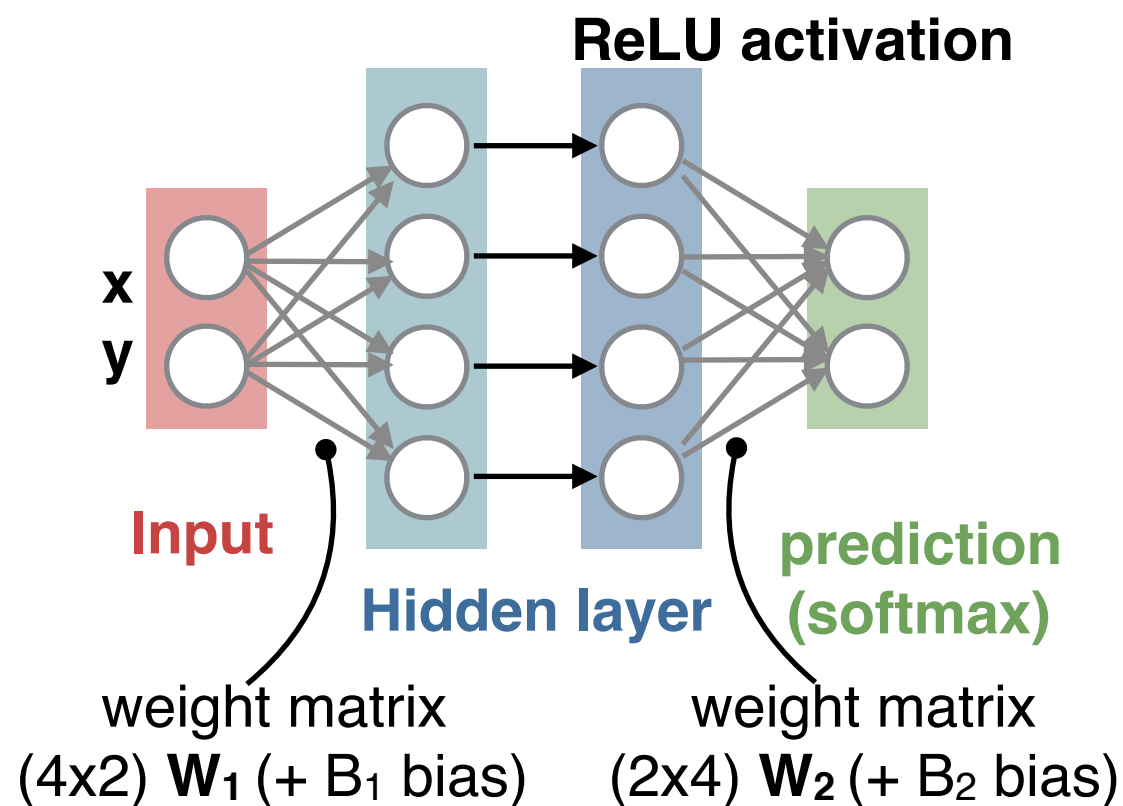
**important: manual tuning of learning rate  $\eta$**

**improved learning with momentum**

$$\vec{w}^{(k)} = \vec{w}^{(k-1)} + \eta \underbrace{(\mu \vec{m}^{(k-1)} - \frac{dL}{d\vec{w}})}_{\text{Momentum } \vec{m}^{(k)}}$$

→ pytorch perform automatic differentiation for you, only forward path definition is necessary

# Multilayer Perceptron




**complex non-linear transfer function**


$$f = W_2 \max(0, W_1 \vec{x} + B_1) + B_2$$

→ differentiable with chain rule (autograd)

**XOR problem**  
→ not linearly separable



→ neural networks (including CNNs) comprise of arbitrarily many hidden layers >100



# Example: Training MLP for XOR classification

pytorch implementation of two-layer MLP with 4 hidden neurons

```
import torch
import torch.nn as nn
import torch.nn.functional as F

#define multilayer perceptron
class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.linear1 = nn.Linear(2,4)
        self.linear2 = nn.Linear(4,2)
    def forward(self, x):
        x = F.relu(self.linear1(x))
        x = self.linear2(x)
        return x

net = MLP()
#loss criterion and SGD optimizer
crossEntropy = nn.CrossEntropyLoss()
```

import necessary torch libraries  
with shorthands nn and F

initial definition has to contain all  
trainable parameters

network operations are only  
defined for forward path (classic  
relu has no trainable weight →

# Example: Training MLP for XOR classification

*#SGD optimizer and training procedure*

```
optimizer = optim.SGD(net.parameters(), lr=0.025, momentum=0.95)
```

```
run_loss = np.zeros(20)
```

*#iterate over several epochs*

```
for epoch in range(20):
```

```
    run_loss[epoch] = 0.0
```

```
    idx_epoch = torch.randperm(32).view(4,8)
```

*#mini-batches consider subsets of whole dataset*

```
    for iter in range(8):
```

```
        idx_iter = idx_epoch[:,iter]
```

```
        optimizer.zero_grad()
```

*#forward path and loss*

```
        input = data[idx_iter, :, :, :]
```

```
        outputs = net(input)
```

```
        loss = crossEntropy(outputs, label[idx_iter, :, :])
```



*#backward path and weight updates*

```
        loss.backward()
```

```
        optimizer.step()
```

```
        run_loss[epoch] += loss.item()
```

**Softmax  
output  
(during  
learning  
iterations)**



**evolution of decision boundary**